

Secure Deduplication Across Files

Nithin V Nath

Advisor: Dr. Bhavana Kanukurthi

Department of Computer Science and Automation
Indian Institute of Science

23-Jun-2016

- 1 Introduction
 - Deduplication
 - Deduplication Example
 - Secure Deduplication
 - Related Work
 - Preliminaries
 - Contributions
 - Road-map
- 2 Construction
 - DD-Across
- 3 Conclusion

Deduplication

- Large amount of data stored in cloud storage.
- Multiple users store the same file.
- Service providers need to employ space saving techniques to keep cost down.

Definition

Technique that enables storage providers to store a single copy of the data.

Deduplication in Action

- Alice uploads a file M to the server S .
- Bob requests to upload his copy of the same file M to S .
- The server identifies that M is already stored and simply updates the metadata associated with M to show that the file is owned by both Alice and Bob.
- Make this an image

Secure Deduplication

- Deduplication along with privacy is a conflicting idea
- Users would like their data to be encrypted
- Storage providers would like to identify the file uploaded by user to enable deduplication.

- Photos taken one after the other are often *almost* identical to each other.
- These multiple files are not supported in traditional file level deduplication.
- **Challenge:** Identify that plaintexts underneath these ciphertexts are close to each other and store only the difference.

Problem Statement



- Convergent Encryption (year?)
- Message Locked Encryption
- Interactive Message Locked Encryption

Convergent Encryption

- Deterministic cryptosystem that produces identical ciphertext files from identical plaintext files.
- $K = H(M)$
- $C = E(K, M)$
- $M = D(K, C)$

Message Locked Encryption

- A cryptographic primitive. $\text{MLE} = (\mathcal{P}, \mathcal{K}, \mathcal{E}, \mathcal{D}, \mathcal{T})$.
- $\mathcal{K}_P(M)$ derives the key from the message.
- \mathcal{T} is the tag generation algorithm.
- Semantic security cannot be achieved using MLE.
- [Include an image]

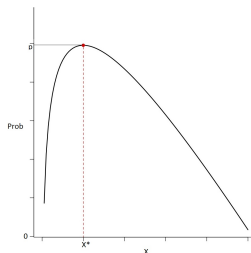
Interactive Message Locked Encryption

- Extension of *MLE*.
- Uses interaction.
- Defined using one algorithm and three protocols
 - 1 $\text{Init}(1^\lambda)$ - The initialization algorithm.
 - 2 Reg - Register a client with the server.
 - 3 $\text{Put}(M, \sigma_C)$ - Puts a plaintext M and returns f , an identifier
 - 4 $\text{Get}(f, \sigma_C)$ - Fetches the file f .

Entropy

- Entropy is a measure of randomness
- Min-entropy of X is the negative log of maximum predictability.

$$H_{\infty}(X) = -\log(\max_x \Pr[X = x])$$



Statistical Distance



Extractors



Deterministic Encryption

- $SE = (E, D)$
- $c \leftarrow E(1^\lambda, k, m)$
- $m \leftarrow D(1^\lambda, k, c)$
- Why is this meaningful in this setting?

Error Correcting Codes

- (\mathcal{M}, K, τ) -code C .
- C is a subset $\{w_0, w_1, \dots, w_K\}$ of \mathcal{M} .
- $\tau > 0$ is the largest number such that there is at most one valid code word $c \in C$ for a message w such that $\text{dis}(w, c) \leq \tau$.
- Enc - The map from i to w_i .
- Dec - The map that finds, given w , the $c \in C$ such that $\text{dis}(w, c) \leq \tau$

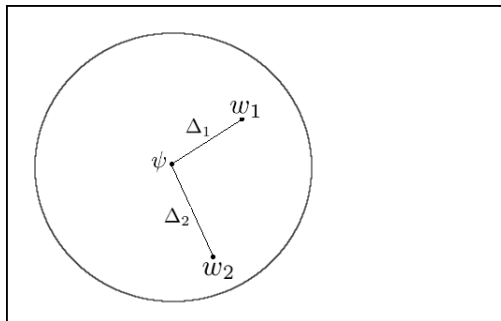
Collision Resistant Hash Functions

- $\mathcal{H} : \{0, 1\}^n \rightarrow \{0, 1\}^m$
- Collision resistant if
 - $m < n$ and
 - for all PPT \mathcal{A} , there exists a negligible function $\text{negl}(\lambda)$ such that for all security parameters $\lambda \in \mathbb{N}$,

$$\Pr[(x_0, x_1) \leftarrow \mathcal{A}(1^\lambda, \mathcal{H}) : x_0 \neq x_1 \wedge \mathcal{H}(x_0) = \mathcal{H}(x_1)] \leq \text{negl}(\lambda)$$

- Family of hash functions: $H = (\mathcal{HK}, \mathcal{H})$

- DD – Across (deduplication across files) which enables deduplication even for files that are close to each other.



- Setting
- Adversarial Model
- Privacy Games
- DD – Across construction
- DD – Across proof

- An honest-but-curious server.
- A set of clients.
- \mathcal{A} can control a subset of these clients.
- Formally modelled using a game G .
- G sets up and controls an instance of a server.

Adversarial Model

- Adversary \mathcal{A} is invoked with access to a set of procedures.
- MSG procedure allows adversary to set up multiple clients and to send arbitrary messages to the server.
- INIT procedure starts protocol instances on behalf of a legitimate client L , using inputs chosen by A .
- STEP procedure advances a protocol instance by running the next step algorithm.
- STATE procedure returns the server's state - including stored ciphertexts, public parameters, etc. Only read only access is gained using this.



- A metric space $(\mathcal{M}, \text{dis})$ with hamming distance as the distance metric.
- An (l, m, κ, ϵ) -strong extractor.
- An error-correcting code $C = (\mathcal{M}, K, \tau)$.
- A collision resistant hash function family $H = (\mathcal{HK}, \mathcal{H})$.
- $SE = (E, D)$ denotes a symmetric encryption scheme.

- DD – Across[C, H, SE].
- Server maintains 3 tables
 - **fil**: which contains the encryptions of the files uploaded by the clients.
 - **delt**: which stores the Δ .
 - **own**: which stores the ownership information.

DD-Across Construction - Init

Init

$S \leftarrow_{\$} \{0, 1\}^{s(\lambda)}$

$K_h \leftarrow_{\$} \mathcal{HK}(1^\lambda)$

$p = (S || K_h)$

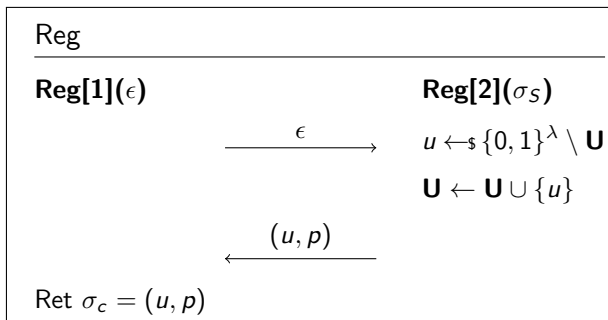
$\mathbf{U} \leftarrow \phi$

$\mathbf{fil} \leftarrow \phi; \mathbf{delt} \leftarrow \phi$

$\mathbf{own} \leftarrow \phi$

Ret $\sigma_S = (p, \mathbf{U}, \mathbf{fil}, \mathbf{delt}, \mathbf{own})$

DD-Across Construction - Reg



DD-Across Construction - Put

Put

Put[1]((u,p),m)

$\psi \leftarrow \text{Dec}(m)$

$k \leftarrow \text{Ext}_\lambda(S, \psi)$

$C_\psi \leftarrow \text{Enc}_{S||K_h}(k, \psi)$

$\Delta \leftarrow \text{Diff}(\psi, m)$

$\xrightarrow{u, C_\psi, \Delta}$

Put[2](σ_S)

$t_1 \leftarrow \mathcal{H}(K_h, C_\psi)$

$t_2 \leftarrow \mathcal{H}(K_h, \Delta)$

$t = (t_1, t_2)$

$\text{SiffE}(\mathbf{fl}, t_1, C_\psi)$

$\text{SiffE}(\mathbf{delt}, t_2, \Delta)$

$\text{SiffE}(\mathbf{own}, (u, t), 1)$

\xleftarrow{t}

Ret (t, k)

DD-Across Construction - Get

Get

Get[1]((u,p),t,k)

Get[2](σ_S)

u, t

$C_\psi \leftarrow \mathbf{fil}[t_1]$

$\Delta \leftarrow \mathbf{delt}[t_2]$

$o \leftarrow \mathbf{own}[u, t]$

if $o = \perp$ **then**

$C_\psi = \perp$

$\Delta = \perp$

C_ψ, Δ

if $C_\psi = \perp$ **then** Ret \perp ; **fi**

$\psi \leftarrow \mathit{Dec}_{S||K_h}(k, C_\psi)$

$m \leftarrow \mathbf{Comb}(\psi, \Delta)$

Ret m

- Recovery is guaranteed.
- For \mathcal{A} to win, a mismatch in the plaintext m put on the server and the plaintext m' recovered using Get.
- Immutability of the tables means once put, it cannot be changed.
- Only possibility is hash collision.

Definition

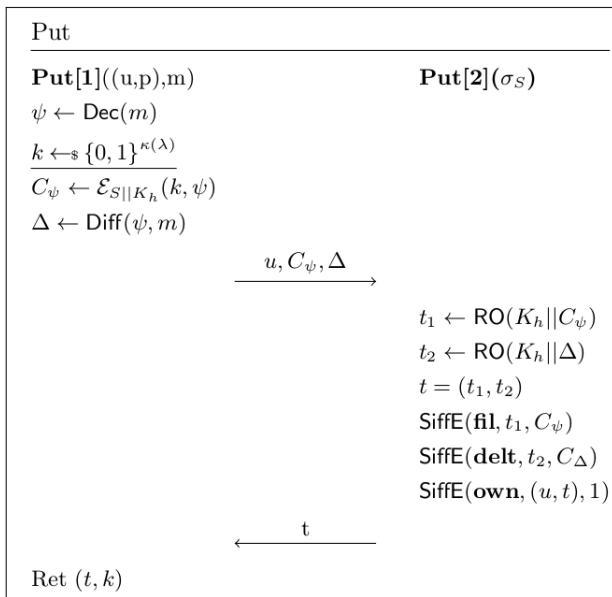
The error-correcting code $C = (\mathcal{M}, K, \tau)$ is said to be compatible with a source S with min-entropy $\mu(\lambda)$ iff $2^{\mu(\lambda) - \tau}$ is negligible.

Theorem

If \mathcal{E} is CPA-secure and the code $C = (\mathcal{M}, K, \tau)$ is compatible with the source S , then $\text{DD} - \text{Across}_{RO}[\mathcal{E}, C]^a$ is PRIV-secure.

^a $\text{DD} - \text{Across}_{RO}$ is the ROM analogue of $\text{DD} - \text{Across}$ which models H as a random oracle

DD-Across Privacy Hybrids



DD-Across Privacy Hybrids

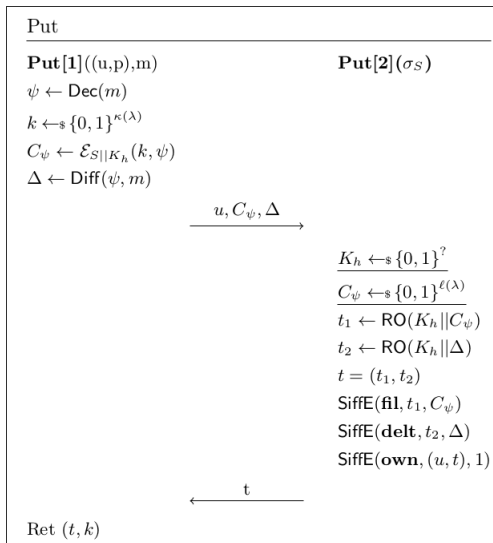


Figure: The Put protocol in game H_3

DD-Across Privacy Hybrids

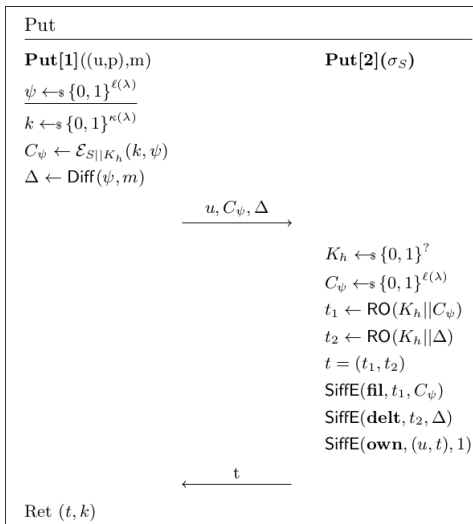


Figure: The Put protocol in game H_4

Summary

- The **first main message** of your talk in one or two lines.
- The **second main message** of your talk in one or two lines.
- Perhaps a **third message**, but not more than that.

Open Problems and Future Work



Block Title

You can also highlight sections of your presentation in a block, with it's own title

Theorem

There are separate environments for theorems, examples, definitions and proofs.

Example

Here is an example of an example block.

For Further Reading I



A. Author.

Handbook of Everything.

Some Press, 1990.



S. Someone.

On this and that.

Journal of This and That, 2(1):50–100, 2000.